

```

from sklearn.preprocessing import LabelBinarizer
from sklearn.metrics import classification_report
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.optimizers import SGD
from tensorflow.keras.datasets import mnist
from tensorflow.keras import backend as K
import matplotlib.pyplot as plt
import numpy as np

((X_train, Y_train), (X_test, Y_test)) = mnist.load_data()
X_train = X_train.reshape((X_train.shape[0], 28 * 28 * 1))
X_test = X_test.reshape((X_test.shape[0], 28 * 28 * 1))
X_train = X_train.astype("float32") / 255.0
X_test = X_test.astype("float32") / 255.0

lb = LabelBinarizer()
Y_train = lb.fit_transform(Y_train)
Y_test = lb.transform(Y_test)

model = Sequential()
model.add(Dense(128, input_shape=(784,), activation="sigmoid"))
model.add(Dense(64, activation="sigmoid"))
model.add(Dense(10, activation="softmax"))

/opt/anaconda3/lib/python3.12/site-
packages/keras/src/layers/core/dense.py:87: UserWarning: Do not pass an
`input_shape`/`input_dim` argument to a layer. When using Sequential models,
prefer using an `Input(shape)` object as the first layer in the model
instead.
  super().__init__(activity_regularizer=activity_regularizer, **kwargs)

sgd = SGD(0.01)
epochs=10
model.compile(loss="categorical_crossentropy",
optimizer=sgd,metrics=["accuracy"])
H = model.fit(X_train, Y_train, validation_data=(X_test, Y_test),epochs=10,
batch_size=128)

Epoch 1/10
469/469 ————— 1s 807us/step - accuracy: 0.1379 - loss: 2.3362
- val_accuracy: 0.2881 - val_loss: 2.2518
Epoch 2/10
469/469 ————— 0s 695us/step - accuracy: 0.3347 - loss: 2.2402
- val_accuracy: 0.5069 - val_loss: 2.1920
Epoch 3/10
469/469 ————— 0s 691us/step - accuracy: 0.4646 - loss: 2.1780
- val_accuracy: 0.4996 - val_loss: 2.1108
Epoch 4/10
469/469 ————— 0s 695us/step - accuracy: 0.5249 - loss: 2.0903

```

```

- val_accuracy: 0.6023 - val_loss: 1.9937
Epoch 5/10
469/469 ————— 0s 703us/step - accuracy: 0.5910 - loss: 1.9660
- val_accuracy: 0.6142 - val_loss: 1.8334
Epoch 6/10
469/469 ————— 0s 696us/step - accuracy: 0.6179 - loss: 1.8009
- val_accuracy: 0.6526 - val_loss: 1.6415
Epoch 7/10
469/469 ————— 0s 712us/step - accuracy: 0.6512 - loss: 1.6115
- val_accuracy: 0.6742 - val_loss: 1.4466
Epoch 8/10
469/469 ————— 0s 709us/step - accuracy: 0.6798 - loss: 1.4220
- val_accuracy: 0.7225 - val_loss: 1.2725
Epoch 9/10
469/469 ————— 0s 762us/step - accuracy: 0.7159 - loss: 1.2537
- val_accuracy: 0.7419 - val_loss: 1.1296
Epoch 10/10
469/469 ————— 0s 726us/step - accuracy: 0.7379 - loss: 1.1217
- val_accuracy: 0.7646 - val_loss: 1.0144

```

```

predictions = model.predict(X_test, batch_size=128)
print(classification_report(Y_test.argmax(axis=1), predictions.argmax(axis=1),
target_names=[str(x) for x in lb.classes_]))

```

```

79/79 ————— 0s 543us/step

```

	precision	recall	f1-score	support
0	0.82	0.96	0.88	980
1	0.78	0.99	0.87	1135
2	0.85	0.76	0.80	1032
3	0.63	0.81	0.71	1010
4	0.71	0.75	0.73	982
5	0.83	0.34	0.49	892
6	0.84	0.86	0.85	958
7	0.80	0.87	0.83	1028
8	0.81	0.60	0.69	974
9	0.67	0.63	0.65	1009
accuracy			0.76	10000
macro avg	0.77	0.76	0.75	10000
weighted avg	0.77	0.76	0.75	10000

```

plt.style.use("ggplot")
plt.figure()
plt.plot(np.arange(0, epochs), H.history["loss"], label="train_loss")
plt.plot(np.arange(0, epochs), H.history["val_loss"], label="val_loss")
plt.plot(np.arange(0, epochs), H.history["accuracy"], label="train_acc")
plt.plot(np.arange(0, epochs), H.history["val_accuracy"], label="val_acc")

```

```
plt.title("Training Loss and Accuracy")
plt.xlabel("Epoch #")
plt.ylabel("Loss/Accuracy")
plt.legend()

<matplotlib.legend.Legend at 0x3141a3170>
```

